

# Lab Assignment 3

20232241 Dookyung Kang

1. In my version of code, the size of moving data is exactly same as the size of vector \* 2. Only remainder of vectors are transferred for the last stream, not aligned to stream size:  $2 * \text{inputLength}$
2. Pinned memory is used when cudaMemcpyAsync called, in order to transfer memory between device and host concurrently, especially using stream.
3. This is entire code.

```
#include <gputk.h>

__global__ void vecAdd(float *in1, float *in2, float *out, int len) {
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    if (index < len) {
        out[index] = in1[index] + in2[index];
    }
}

#ifdef STREAM
#define stream 4
#endif

int main(int argc, char **argv) {
    gpuTKArg_t args;
    int inputLength;
    float *hostInput1;
    float *hostInput2;
    float *hostOutput;
    float *deviceInput1;
    float *deviceInput2;
    float *deviceOutput;
    unsigned int numStreams;

    args = gpuTKArg_read(argc, argv);

    gpuTKTime_start(Generic, "Importing data and creating memory on host");
    hostInput1 =
        (float *)gpuTKImport(gpuTKArg_getInputFile(args, 0), &inputLength);
    hostInput2 =
        (float *)gpuTKImport(gpuTKArg_getInputFile(args, 1), &inputLength);
    hostOutput = (float *)malloc(inputLength * sizeof(float));
    gpuTKTime_stop(Generic, "Importing data and creating memory on host");

    gpuTKLog	TRACE, "The input length is ", inputLength);

    gpuTKTime_start(GPU, "Allocating Pinned memory.");

    ///@@ Allocate GPU memory here using pinned memory here
    cudaMallocHost((void **)&deviceInput1, inputLength * sizeof(float));
    cudaMallocHost((void **)&deviceInput2, inputLength * sizeof(float));
    cudaMallocHost((void **)&deviceOutput, inputLength * sizeof(float));

    ///@@ Create and setup streams
    numStreams = STREAM;
    cudaStream_t streams[numStreams];
    for (int i = 0; i < numStreams; i++) {
        cudaStreamCreate(&streams[i]);
    }

    ///@@ Calculate data segment size of input data processed by each stream
    int streamSizes[numStreams];
    int offsets[numStreams];
```

```

int streamSizeBase = inputLength / numStreams;
for (int i = 0; i < numStreams; i++) {
    if (i < numStreams - 1)
        streamSizes[i] = streamSizeBase;
    else
        streamSizes[i] = inputLength - streamSizeBase * (numStreams - 1);
    offsets[i] = i * streamSizeBase;
}

int blockSize = 256;
int numBlocks = (inputLength + blockSize - 1) / blockSize;

gpuTKTime_start(Compute, "Performing CUDA computation");
///< Perform parallel vector addition with different streams.
for (unsigned int s = 0; s < numStreams; s++){
    ///< Asynchronous copy data to the device memory in segments
    ///< Calculate starting and ending indices for per-stream data
    int offset = offsets[s];
    int streamSize = streamSizes[s];
    cudaStream_t stream = streams[s];

    cudaMemcpyAsync(&deviceInput1[offset], &hostInput1[offset],
        streamSize * sizeof(float), cudaMemcpyHostToDevice, stream);
    cudaMemcpyAsync(&deviceInput2[offset], &hostInput2[offset],
        streamSize * sizeof(float), cudaMemcpyHostToDevice, stream);

    ///< Invoke CUDA Kernel
    ///< Determine grid and thread block sizes (consider occupancy)
    vecAdd<<<numBlocks, blockSize, 0, stream>>>
        (&deviceInput1[offset], &deviceInput2[offset], &deviceOutput[offset], streamSize);

    ///< Asynchronous copy data from the device memory in segments
    cudaMemcpyAsync(&hostOutput[offset], &deviceOutput[offset],
        streamSize * sizeof(float), cudaMemcpyDeviceToHost, stream);
}

///< Synchronize
for (int i = 0; i < numStreams; i++) {
    cudaStreamSynchronize(streams[i]);
}

gpuTKTime_stop(Compute, "Performing CUDA computation");

gpuTKTime_start(GPU, "Freeing Pinned Memory");
///< Destroy cudaStream
for (int i = 0; i < numStreams; i++) {
    cudaStreamDestroy(streams[i]);
}

///< Free the GPU memory here
cudaFreeHost(deviceInput1);
cudaFreeHost(deviceInput2);
cudaFreeHost(deviceOutput);

gpuTKTime_stop(GPU, "Freeing Pinned Memory");

gpuTKSolution(args, hostOutput, inputLength);

free(hostInput1);
free(hostInput2);
free(hostOutput);

return 0;
}

```

4. Evaluation is performed in NVIDIA RTX A5000, with CUDA 12.0 environment.

N	16	64	93	112	1120	9921	14000	25365	48000	96000
<b>Allocating (ms)</b>	0.20638	0.438018	0.594588	0.687433	5.68884	49.6007	69.8475	126.265	239.515	477.572
<b>Computation (ms)</b>	0.106404	0.108997	0.108298	0.104469	0.108586	0.125712	0.138378	0.16151	0.187456	0.264895
<b>Freeing (ms)</b>	0.348197	0.34155	0.341255	0.343108	0.340594	0.350368	0.34517	0.342886	0.357332	0.36363

5. Evaluation is performed in NVIDIA RTX A5000, with CUDA 12.0 environment.

N	2	4	8	12	16	24	32
<b>Allocating (ms)</b>	238.669	239.327	238.956	239.956	238.905	239.025	238.824
<b>Computation (ms)</b>	0.144958	0.181115	0.244066	0.320699	0.36287	0.467215	0.562694
<b>Freeing (ms)</b>	0.364924	0.353059	0.355825	0.367033	0.366532	0.38028	0.391648